

Heterogeneous Computing for AI - Lecture ~03

Advanced Concurrency in Python

Raghava Mukkamala

Associate Professor, Director for Centre for Business Data Analytics (bda.cbs.dk)

Department of Digitalization, Copenhagen Business School, Denmark

Email: rrm.digi@cbs.dk

Associate Professor, Department of Technology,
Kristiania University College, Oslo, Norway

Many slides are taken from the following authors with due respect to their contributions.

1) An Introduction to Python Concurrency by David Beazley

Outline

- Recap - Threading
- Threading and Race Conditions
- Thread Synchronization Primitives
- Mutex Locks
- Semaphores
- Events and Condition Variables
- Thread-safe data structures: Queues

Learning goals for today

Theoretical

- Gain knowledge about various thread synchronization primitives in Python
- Learn the foundations of basic primitives for concurrency in Python
- Understand the foundations of advanced primitives for concurrency in Python

Practical

- Be able to program using basic thread synchronization primitives
- Understand how to use advanced thread synchronization primitives

Recap - Threading

Python Thread Programming

Thread Basics

`% python program.py`

`statement`
`statement`

`...`

`create thread(foo)`

`statement`
`statement`

`...`

`statement`
`statement`

`...`



Key idea: Thread is like a little "task" that independently runs inside your program

thread

`def foo():`
 `statement`
 `statement`
 `...`
 `return or exit`

Python Thread Programming

Joining a Thread

- Once you start a thread, it runs independently
- Use `t.join()` to wait for a thread to exit

```
t.start()           # Launch a thread
...
# Do other work
...
# Wait for thread to finish
t.join()            # Waits for thread t to exit
```

- This only works from *other* threads
- A thread can't join itself

Python Thread Programming

Daemonic Threads

- If a thread runs forever, make it "daemonic"

```
t.daemon = True  
t.setDaemon(True)
```

- If you don't do this, the interpreter will lock when the main thread exits---waiting for the thread to terminate (which never happens)
- Normally you use this for background tasks

Threading and Race Conditions

Threading and Race Conditions

Interlude

- Creating threads is really easy
- You can create thousands of them if you want
- Programming with threads is hard
- Really hard

Q: Why did the multithreaded chicken cross the road?

A: to To other side. get the

-- Jason Whittington

Threading and Race Conditions

Access to Shared Data

- Threads share all of the data in your program
- Thread scheduling is non-deterministic
- Operations often take several steps and might be interrupted mid-stream (non-atomic)
- Thus, access to any kind of shared data is also non-deterministic (which is a really good way to have your head explode)

Threading and Race Conditions

Accessing Shared Data

- Consider a shared object

```
x = 0
```

- And two threads that modify it

```
Thread-1  
-----  
...  
x = x + 1  
...
```

```
Thread-2  
-----  
...  
x = x - 1  
...
```

- It's possible that the resulting value will be unpredictably corrupted

Threading and Race Conditions

Accessing Shared Data

- The two threads

Thread-1

...
x = x + 1
...

Thread-2

...
x = x - 1
...

- Low level interpreter execution

Thread-1

↓
LOAD_GLOBAL 1 (x)
LOAD_CONST 2 (1)

BINARY_ADD
STORE_GLOBAL 1 (x)

Thread-2

LOAD_GLOBAL 1 (x)
LOAD_CONST 2 (1)
BINARY_SUB
STORE_GLOBAL 1 (x)

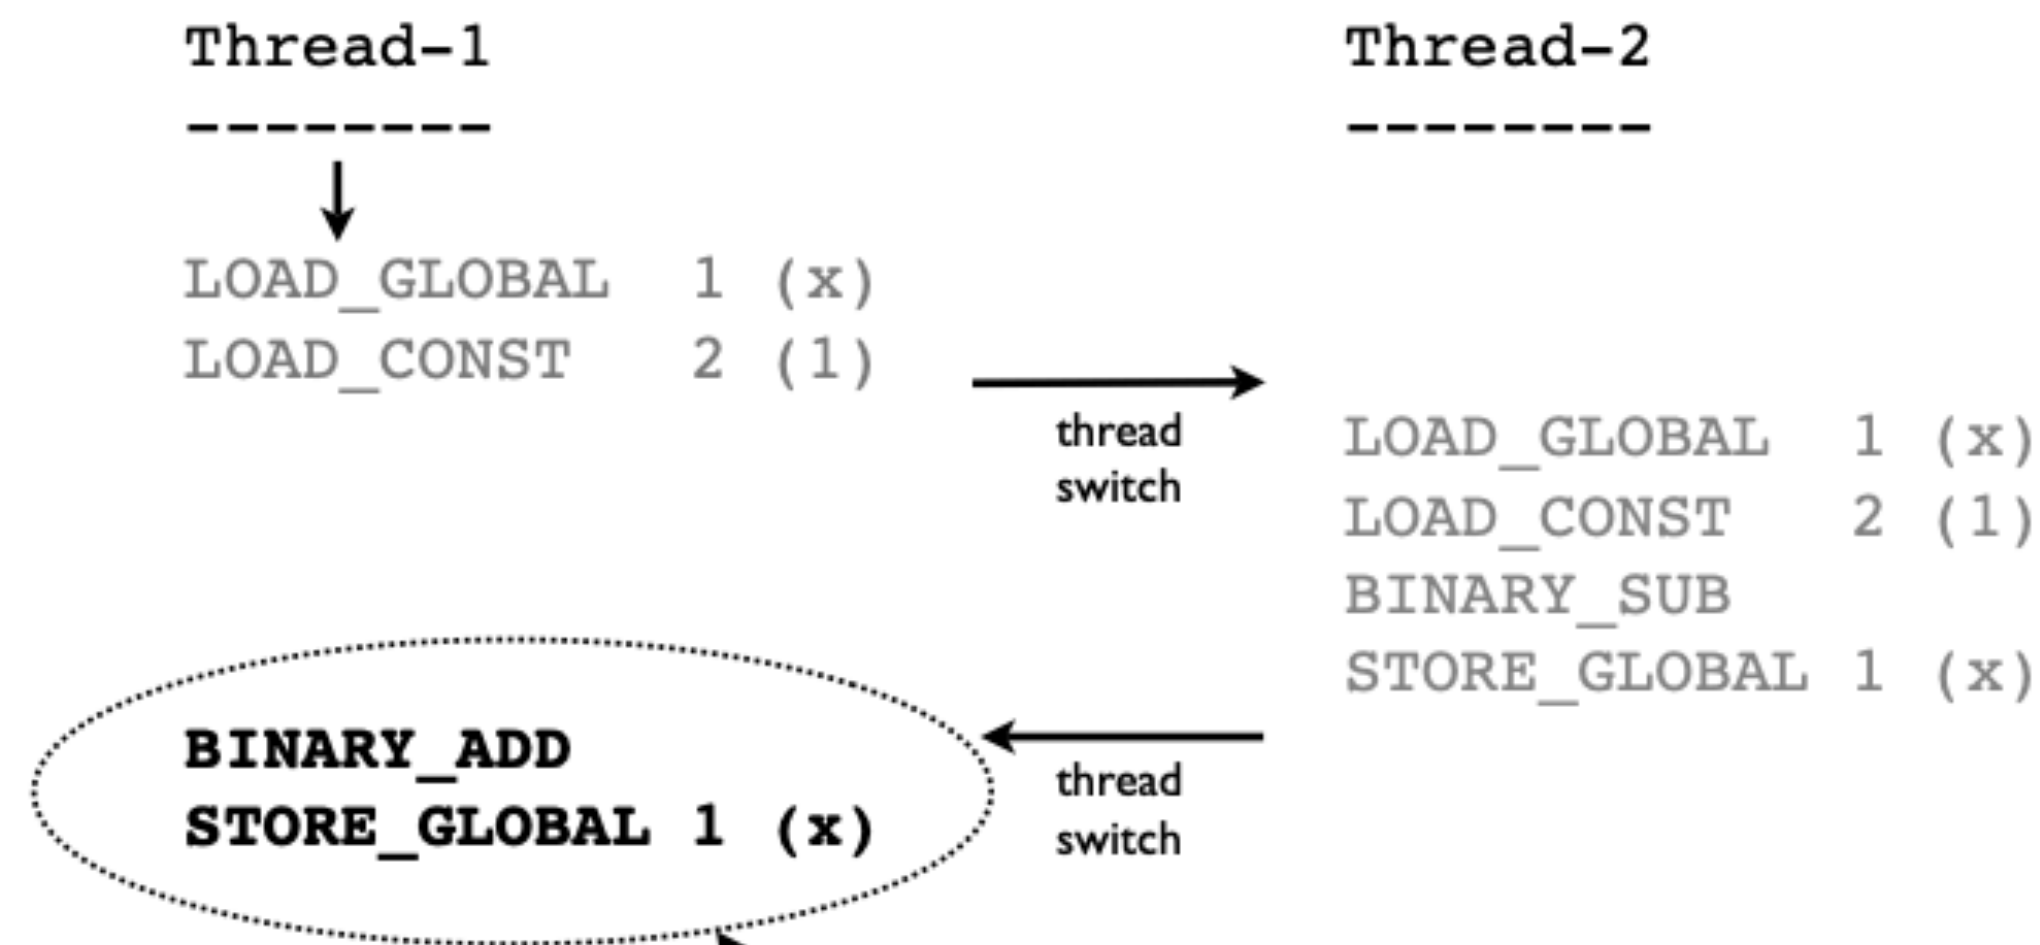
→
thread
switch

←
thread
switch

Threading and Race Conditions

Accessing Shared Data

- Low level interpreter code

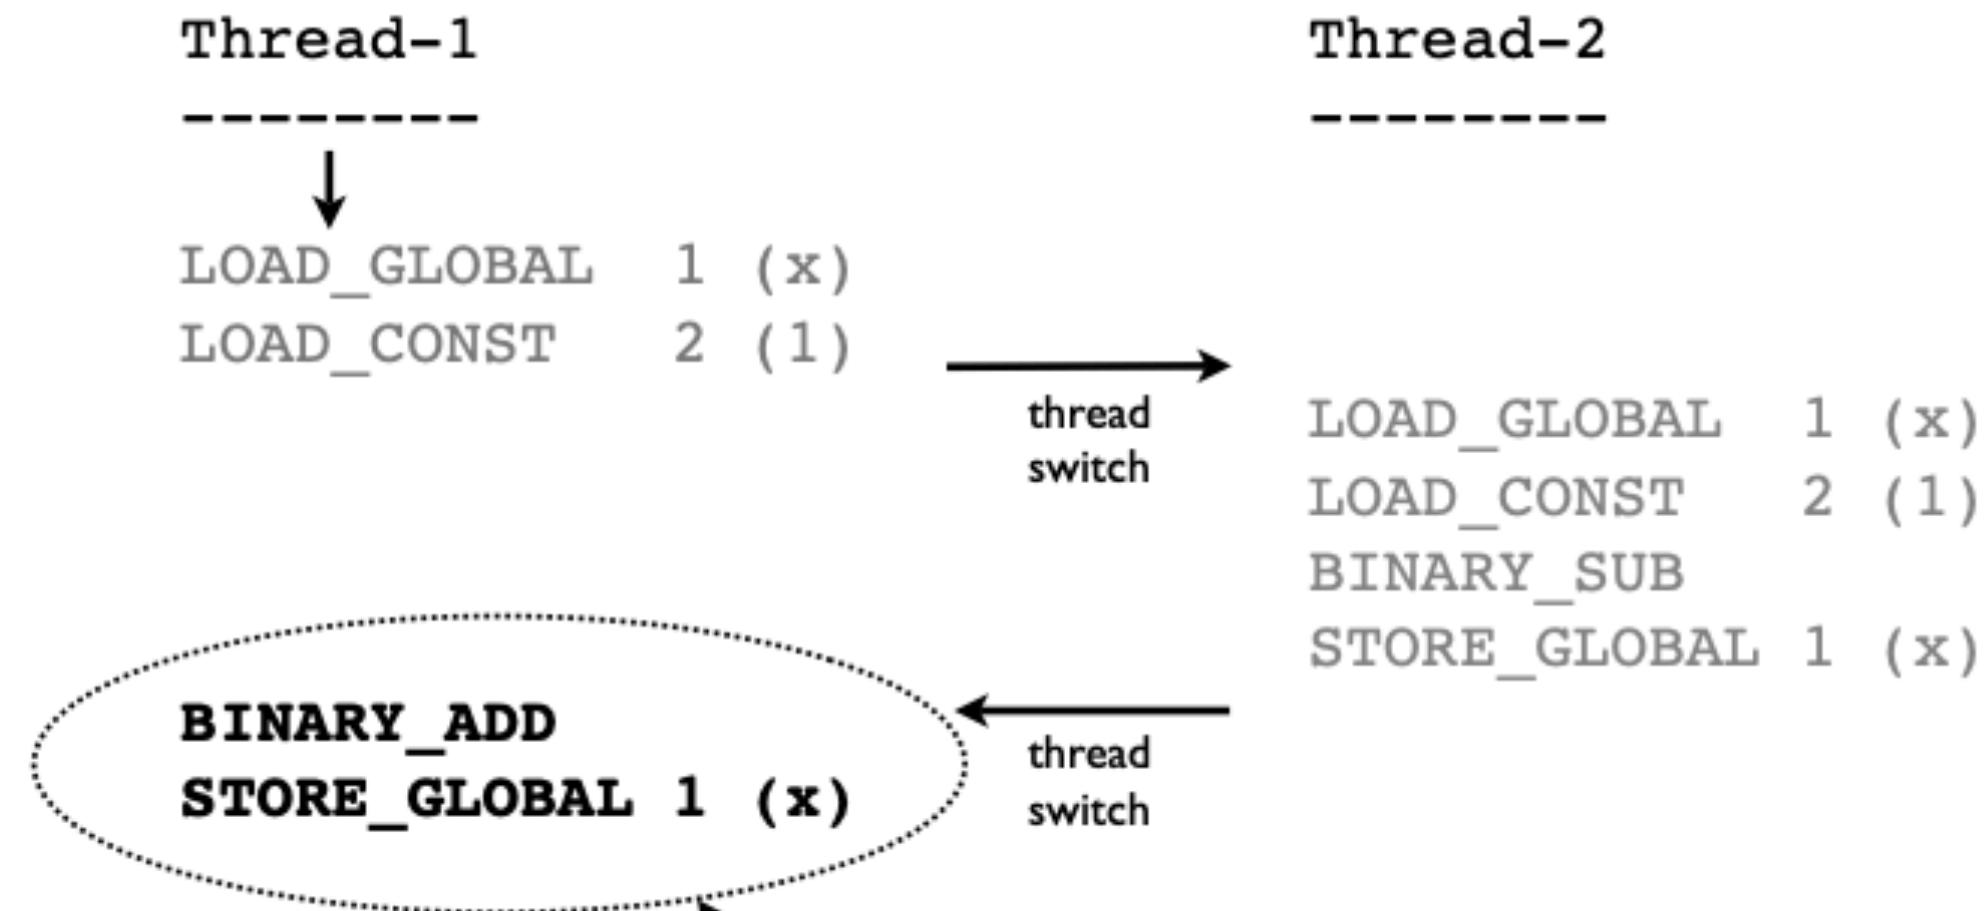


These operations get performed with a "stale" value of x. The computation in Thread-2 is lost.

Threading and Race Conditions

Accessing Shared Data

- Low level interpreter code



These operations get performed with a "stale" value of x. The computation in Thread-2 is lost.

Threading and Race Conditions

Accessing Shared Data

- Is this actually a real concern?

```
x = 0                # A shared value
def foo():
    global x
    for i in xrange(100000000): x += 1

def bar():
    global x
    for i in xrange(100000000): x -= 1

t1 = threading.Thread(target=foo)
t2 = threading.Thread(target=bar)
t1.start(); t2.start()
t1.join(); t2.join()  # Wait for completion
print x              # Expected result is 0
```

- Yes, the print produces a random nonsensical value each time (e.g., -83412 or 1627732)

example in
tutorial notebook

Threading and Race Conditions

Race Conditions

- The corruption of shared data due to thread scheduling is often known as a "race condition."
- It's often quite diabolical--a program may produce slightly different results each time it runs (even though you aren't using any random numbers)
- Or it may just flake out mysteriously once every two weeks

Threading and Race Conditions

Thread Synchronization

- Identifying and fixing a race condition will make you a better programmer (e.g., it "builds character")
- However, you'll probably never get that month of your life back...
- To fix : You have to synchronize threads

Thread Synchronization Primitives

Threading and Race Conditions

Synchronization Options

- The threading library defines the following objects for synchronizing threads
 - Lock
 - RLock
 - Semaphore
 - BoundedSemaphore
 - Event
 - Condition

Mutex Locks

Mutex Locks

Mutex Locks

- Mutual Exclusion Lock

```
m = threading.Lock()
```

- Probably the most commonly used synchronization primitive
- Primarily used to synchronize threads so that only one thread can make modifications to shared data at any given time

Mutex Locks

Mutex Locks

- There are two basic operations

<code>m.acquire()</code>	<code># Acquire the lock</code>
<code>m.release()</code>	<code># Release the lock</code>

- Only one thread can successfully acquire the lock at any given time
- If another thread tries to acquire the lock when its already in use, it gets blocked until the lock is released

Mutex Locks

Use of Mutex Locks

- Commonly used to enclose critical sections

```
x = 0
x_lock = threading.Lock()
```

	Thread-1	Thread-2
	-----	-----

	x_lock.acquire()	x_lock.acquire()
Critical Section	x = x + 1	x = x - 1
	x_lock.release()	x_lock.release()

- Only one thread can execute in critical section at a time (lock gives exclusive access)

example in
tutorial notebook

Mutex Locks

Using a Mutex Lock

- It is your responsibility to identify and lock all "critical sections"

```
x = 0
x_lock = threading.Lock()
```

Thread-1

```
...
x_lock.acquire()
x = x + 1
x_lock.release()
...
```

Thread-2

```
...
x = x - 1
...
```

↑

If you use a lock in one place, but not another, then you're missing the whole point. All modifications to shared state must be enclosed by lock acquire()/release().

Mutex Locks

Locking Perils

- Locking looks straightforward
- Until you start adding it to your code
- Managing locks is a lot harder than it looks

Mutex Locks

Lock Management

- Acquired locks must always be released
- However, it gets evil with exceptions and other non-linear forms of control-flow
- Always try to follow this prototype:

```
x = 0
x_lock = threading.Lock()

# Example critical section
x_lock.acquire()
try:
    statements using x
finally:
    x_lock.release()
```

Mutex Locks

Lock Management

- Python 2.6/3.0 has an improved mechanism for dealing with locks and critical sections

```
x = 0
x_lock = threading.Lock()

# Critical section
with x_lock:
    statements using x
...
```

- This automatically acquires the lock and releases it when control enters/exits the associated block of statements

Mutex Locks

Locks and Deadlock

- Don't write code that acquires more than one mutex lock at a time

```
x = 0
y = 0
x_lock = threading.Lock()
y_lock = threading.Lock()

with x_lock:
    statements using x
    ...
    with y_lock:
        statements using x and y
        ...
```

- This almost invariably ends up creating a program that mysteriously deadlocks (even more fun to debug than a race condition)

Semaphores

Semaphores

Semaphores

- A counter-based synchronization primitive

```
m = threading.Semaphore(n) # Create a semaphore
m.acquire()                # Acquire
m.release()                # Release
```

- `acquire()` - Waits if the count is 0, otherwise decrements the count and continues
- `release()` - Increments the count and signals waiting threads (if any)
- Unlike locks, `acquire()/release()` can be called in any order and by any thread

Semaphores

Semaphore Uses

- Resource control. You can limit the number of threads performing certain operations. For example, performing database queries, making network connections, etc.
- Signaling. Semaphores can be used to send "signals" between threads. For example, having one thread wake up another thread.

Semaphores

Resource Control

- Using a semaphore to limit resources

```
sema = threading.Semaphore(5)    # Max: 5-threads
```

```
def fetch_page(url):  
    sema.acquire()  
    try:  
        u = urllib.urlopen(url)  
        return u.read()  
    finally:  
        sema.release()
```

- In this example, only 5 threads can be executing the function at once (if there are more, they will have to wait)

Semaphores

Thread Signaling

- Using a semaphore to signal


```
done = threading.Semaphore(0)
```

Thread 1

```
...  
statements  
statements  
statements  
done.release()
```

Thread 2

```
done.acquire()  
statements  
statements  
statements  
...
```



- Here, `acquire()` and `release()` occur in different threads and in a different order
- Often used with producer-consumer problems

example in
tutorial notebook

Events and Condition Variables

Events

Events

- Event Objects

```
e = threading.Event()  
e.isSet()      # Return True if event set  
e.set()        # Set event  
e.clear()      # Clear event  
e.wait()       # Wait for event
```

- This can be used to have one or more threads wait for something to occur
- Setting an event will unblock all waiting threads simultaneously (if any)
- Common use : barriers, notification

Events

Event Example

- Using an event to ensure proper initialization

```
init = threading.Event()

def worker():
    init.wait()      # Wait until initialized
    statements
    ...

def initialize():
    statements      # Setting up
    statements      # ...
    ...
    init.set()      # Done initializing

Thread(target=worker).start()  # Launch workers
Thread(target=worker).start()
Thread(target=worker).start()
initialize()                  # Initialize
```

Copyright (C) 2009, David Beazley, <http://www.dabeaz.com>

64

example in
tutorial notebook

Event Example

- Using an event to signal "completion"

```
def master():
    ...
    item = create_item()
    evt = Event()
    worker.send((item, evt))
    ...
    # Other processing
    ...
    ...
    ...
    # Wait for worker
    evt.wait()
```

Worker Thread

```
item, evt = get_work()
processing
processing
...
...
# Done
evt.set()
```

- Might use for asynchronous processing, etc.

Copyright (C) 2009, David Beazley, <http://www.dabeaz.com>

65

Condition Variables

Condition Variables

- Condition Objects

```
cv = threading.Condition([lock])
cv.acquire()      # Acquire the underlying lock
cv.release()      # Release the underlying lock
cv.wait()         # Wait for condition
cv.notify()       # Signal that a condition holds
cv.notifyAll()    # Signal all threads waiting
```

- A combination of locking/signaling
- Lock is used to protect code that establishes some sort of "condition" (e.g., data available)
- Signal is used to notify other threads that a "condition" has changed state

Condition Variables

Condition Variables

- Common Use : Producer/Consumer patterns

```
items = []  
items_cv = threading.Condition()
```

Producer Thread

```
item = produce_item()  
with items_cv:  
    items.append(item)
```

Consumer Thread

```
with items_cv:  
    ...  
    x = items.pop(0)  
  
    # Do something with x  
    ...
```

- First, you use the locking part of a CV
synchronize access to shared data (items)

Condition Variables

Condition Variables

- Common Use : Producer/Consumer patterns


```
items = []  
items_cv = threading.Condition()
```

Producer Thread

```
item = produce_item()  
with items_cv:  
    items.append(item)  
    items_cv.notify()
```

Consumer Thread

```
with items_cv:  
    while not items:  
        items_cv.wait()  
    x = items.pop(0)  
  
    # Do something with x  
    ...
```



- Next you add signaling and waiting
- Here, the producer signals the consumer that it put data into the shared list

Condition Variables

Condition Variables

- Some tricky bits involving wait()
- Before waiting, you have to acquire the lock
- wait() releases the lock when waiting and reacquires when woken
- Conditions are often transient and may not hold by the time wait() returns. So, you must always double-check (hence, the while loop)

Consumer Thread

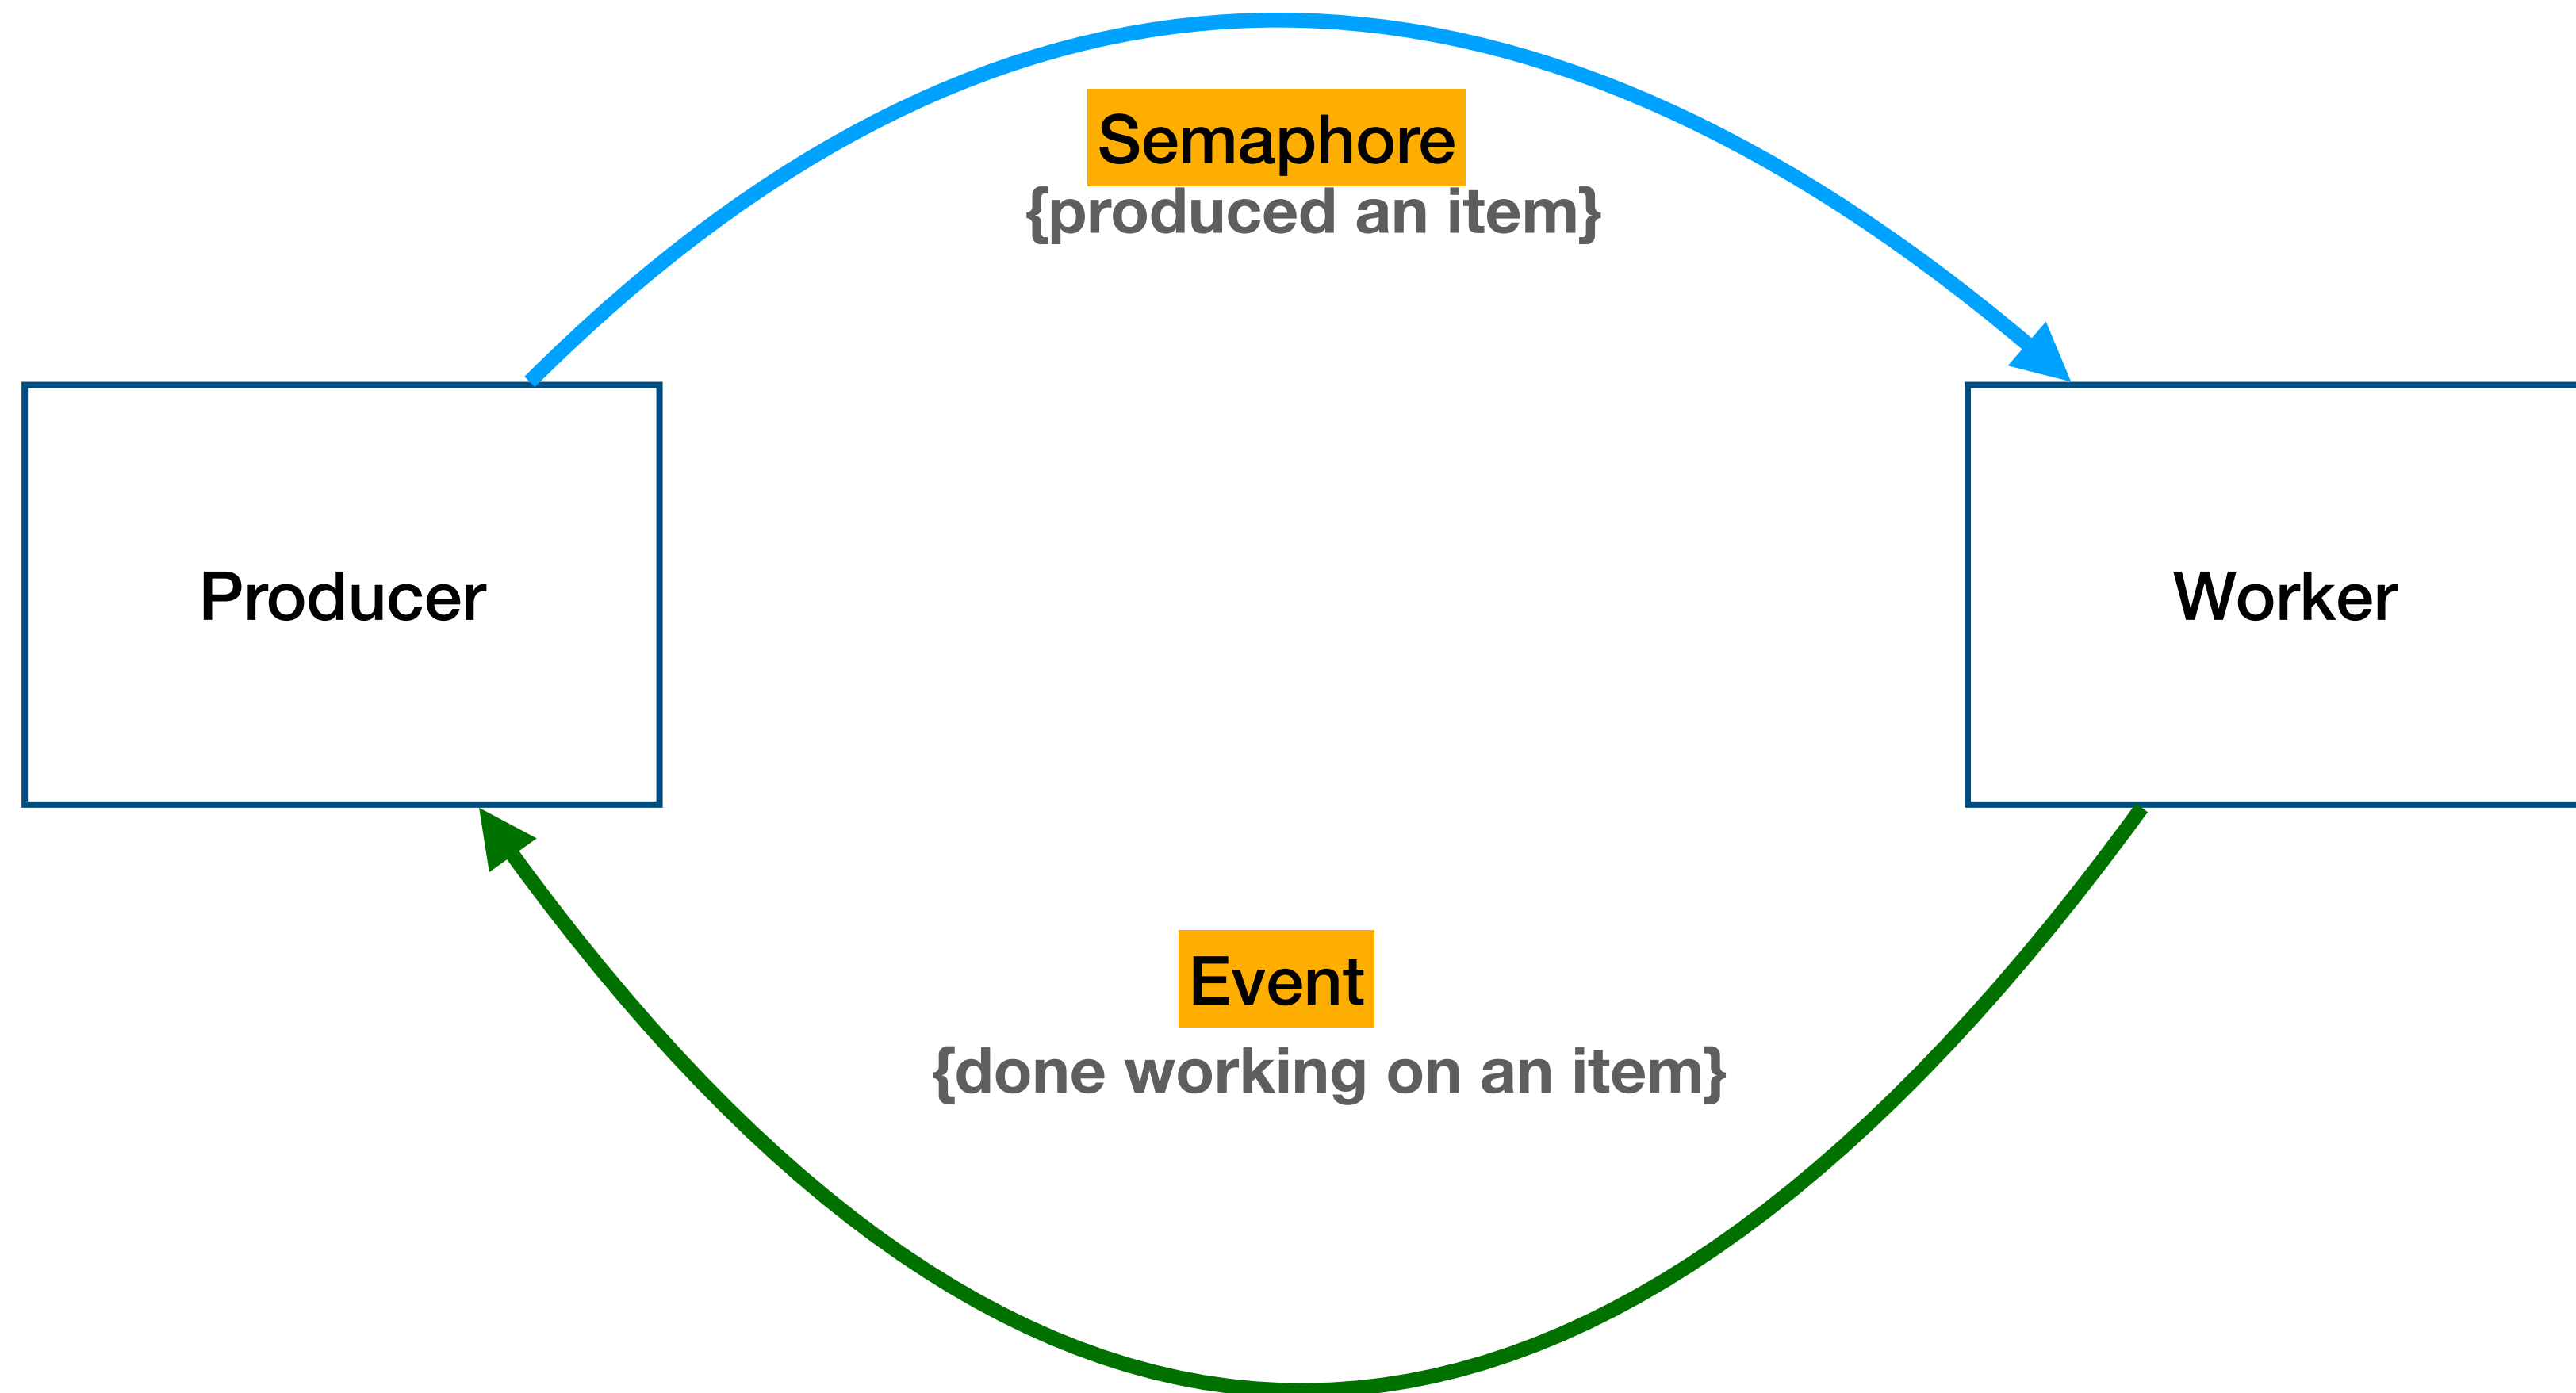
```
with items_cv:  
    while not items:  
        items_cv.wait()  
    x = items.pop(0)  
  
    # Do something with x  
    ...
```


Condition Variables

Interlude

- Working with all of the synchronization primitives is a lot trickier than it looks
- There are a lot of nasty corner cases and horrible things that can go wrong
- Bad performance, deadlock, livelock, starvation, bizarre CPU scheduling, etc...
- All are valid reasons to not use threads

Example on Events and Semaphore




Threads and Queues

Threads and Queues

Threads and Queues

- Threaded programs are often easier to manage if they can be organized into producer/consumer components connected by queues



- Instead of "sharing" data, threads only coordinate by sending data to each other
- Think Unix "pipes" if you will...

Threads and Queues

Queue Library Module

- Python has a thread-safe queuing module
- Basic operations

```
from Queue import Queue
```

```
q = Queue([maxsize])    # Create a queue
q.put(item)              # Put an item on the queue
q.get()                  # Get an item from the queue
q.empty()                # Check if empty
q.full()                 # Check if full
```

- Usage : You try to strictly adhere to get/put operations. If you do this, you don't need to use other synchronization primitives.

Threads and Queues

Queue Usage

- Most commonly used to set up various forms of producer/consumer problems

```
from Queue import Queue  
q = Queue()
```

Producer Thread

```
for item in produce_items():  
    q.put(item)
```

Consumer Thread

```
while True:  
    item = q.get()  
    consume_item(item)
```

- Critical point : You don't need locks here

Threads and Queues

Queue Signaling

- Queues also have a signaling mechanism

```
q.task_done()      # Signal that work is done
q.join()           # Wait for all work to be done
```

- Many Python programmers don't know about this (since it's relatively new)
- Used to determine when processing is done

Producer Thread

```
for item in produce_items():
    q.put(item)
# Wait for consumer
q.join()
```


Consumer Thread

```
while True:
    item = q.get()
    consume_item(item)
    q.task_done()
```

Threads and Queues

Queue Programming

- There are many ways to use queues
- You can have as many consumers/producers as you want hooked up to the same queue



- In practice, try to keep it simple

Resources

- **An Introduction to Python Concurrency.** <http://www.dabeaz.com/usenix2009/concurrent/index.html>
- **Python threads synchronization: Locks, RLocks, Semaphores, Conditions and Queues** <http://www.laurentluce.com/posts/python-threads-synchronization-locks-rlocks-semaphores-conditions-events-and-queues/>
- **Multithreading in Python | Set 1** <https://www.geeksforgeeks.org/multithreading-python-set-1/>
- **Multithreading in Python | Set 2 (Synchronization)** <https://www.geeksforgeeks.org/multithreading-in-python-set-2-synchronization/>

The End