

# Heterogeneous Computing for AI - Lecture ~04

## Parallel Programming using Multiple Processors in Python

**Raghava Mukkamala**

*Associate Professor*, Director for Centre for Business Data Analytics ([bda.cbs.dk](http://bda.cbs.dk))

Department of Digitalization, Copenhagen Business School, Denmark

Email: [rrm.digi@cbs.dk](mailto:rrm.digi@cbs.dk)

Associate Professor, Department of Technology,  
Kristiania University College, Oslo, Norway

Many slides are taken from the following authors with due respect to their contributions.

1) An Introduction to Python Concurrency by David Beazley

# Outline

- The Inside Story on Python Threads
- Processes and Messages
- Multiprocessing Module
- Multiprocessing Module
  - Pipes
  - Queues
  - Process Pools

# Learning goals for today

## Theoretical

- Gain knowledge about various multi-process parallel programming primitives in Python
- Understand the foundations of Process Pools, Queues, Pipes, and others

## Practical

- Be able to write programs that can run on multicore processors/CPU's
- Understand how to pass messages between different processes

# **The Inside Story on Python Threads**

# Python and GIL Lock


## The Infamous GIL

- Here's the rub...
- Only one Python thread can execute in the interpreter at once
- There is a "global interpreter lock" that carefully controls thread execution
- The GIL ensures that sure each thread gets exclusive access to the entire interpreter internals when it's running

# Python and GIL Lock

## GIL Behavior

- Whenever a thread runs, it holds the GIL
- However, the GIL is released on blocking I/O




- So, any time a thread is forced to wait, other "ready" threads get their chance to run
- Basically a kind of "cooperative" multitasking

# Python and GIL Lock

## CPU Bound Processing

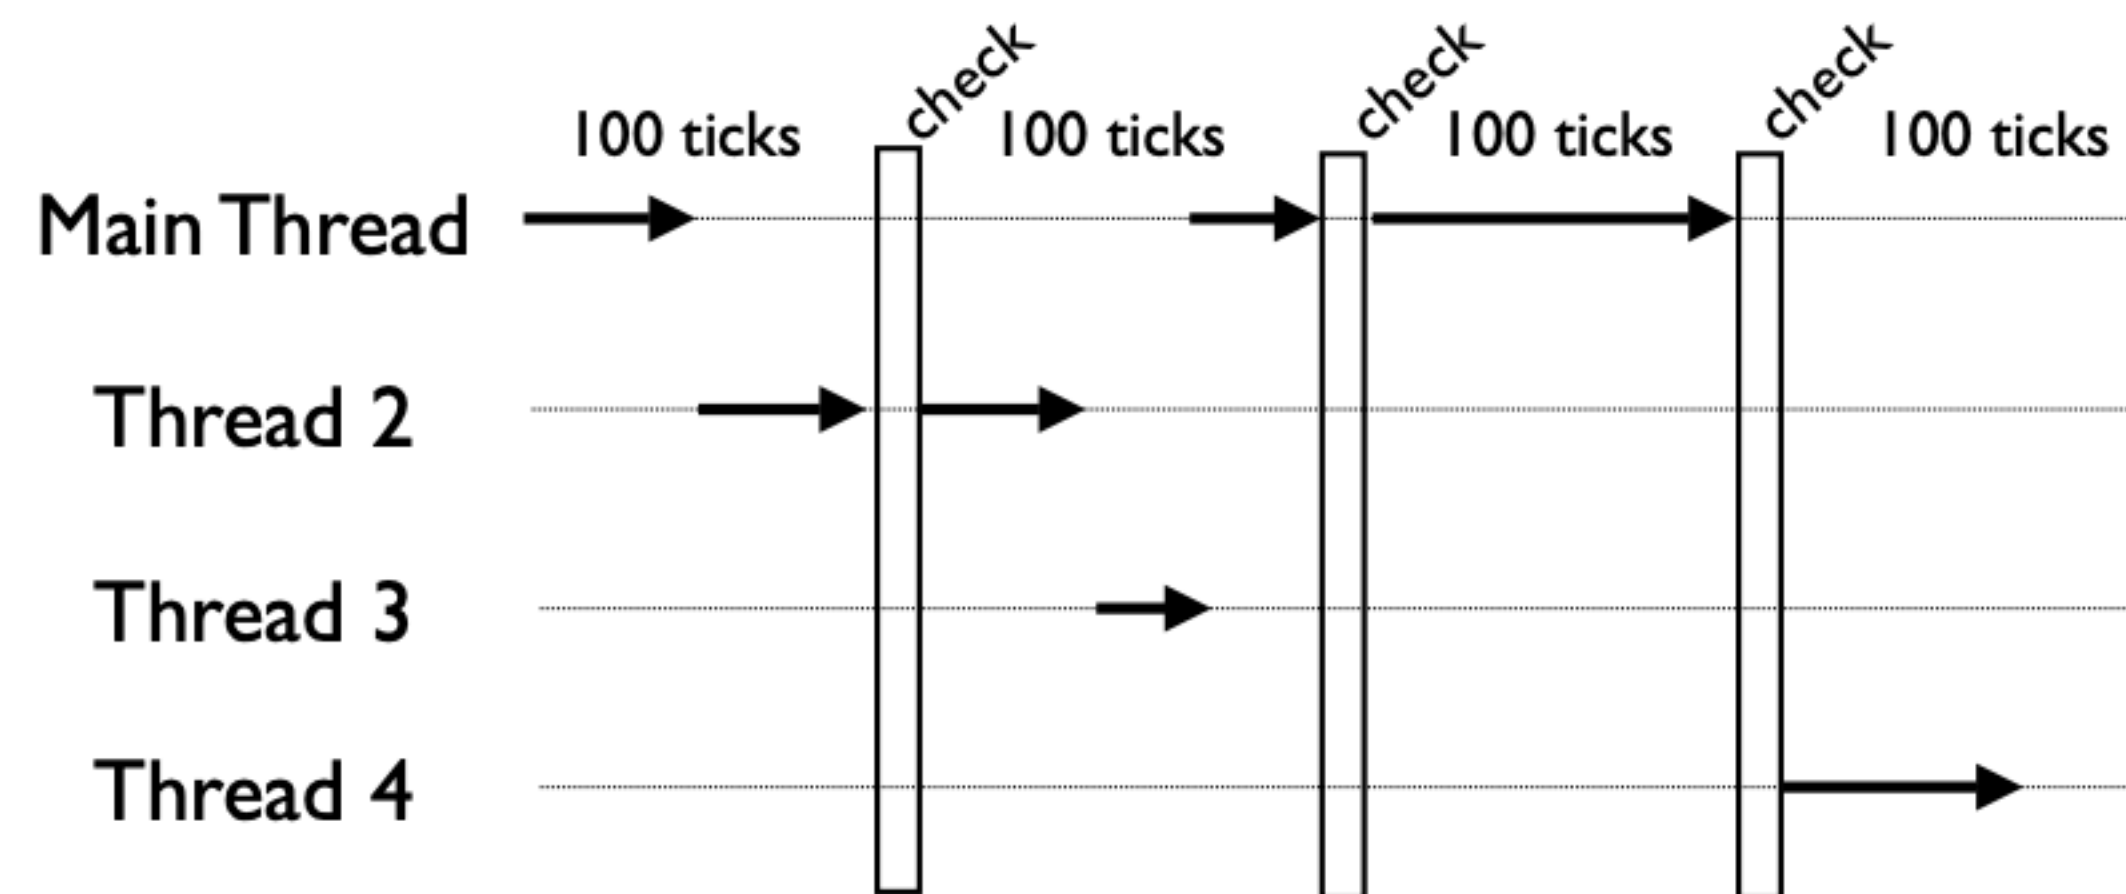
- To deal with CPU-bound threads, the interpreter periodically performs a "check"
- By default, every 100 interpreter "ticks"



# Python and GIL Lock

## The Check Interval

- The check interval is a global counter that is completely independent of thread scheduling



- A "check" is simply made every 100 "ticks"



# Python and GIL Lock

## The Periodic Check

- What happens during the periodic check?
  - In the main thread only, signal handlers will execute if there are any pending signals
  - Release and reacquisition of the GIL
- That last bullet describes how multiple CPU-bound threads get to run (by briefly releasing the GIL, other threads get a chance to run).

# Python and GIL Lock

## Thread Scheduling

- Python does not have a thread scheduler
- There is no notion of thread priorities, preemption, round-robin scheduling, etc.
- For example, the list of threads in the interpreter isn't used for anything related to thread execution
- All thread scheduling is left to the host operating system (e.g., Linux, Windows, etc.)

# Python and GIL Lock

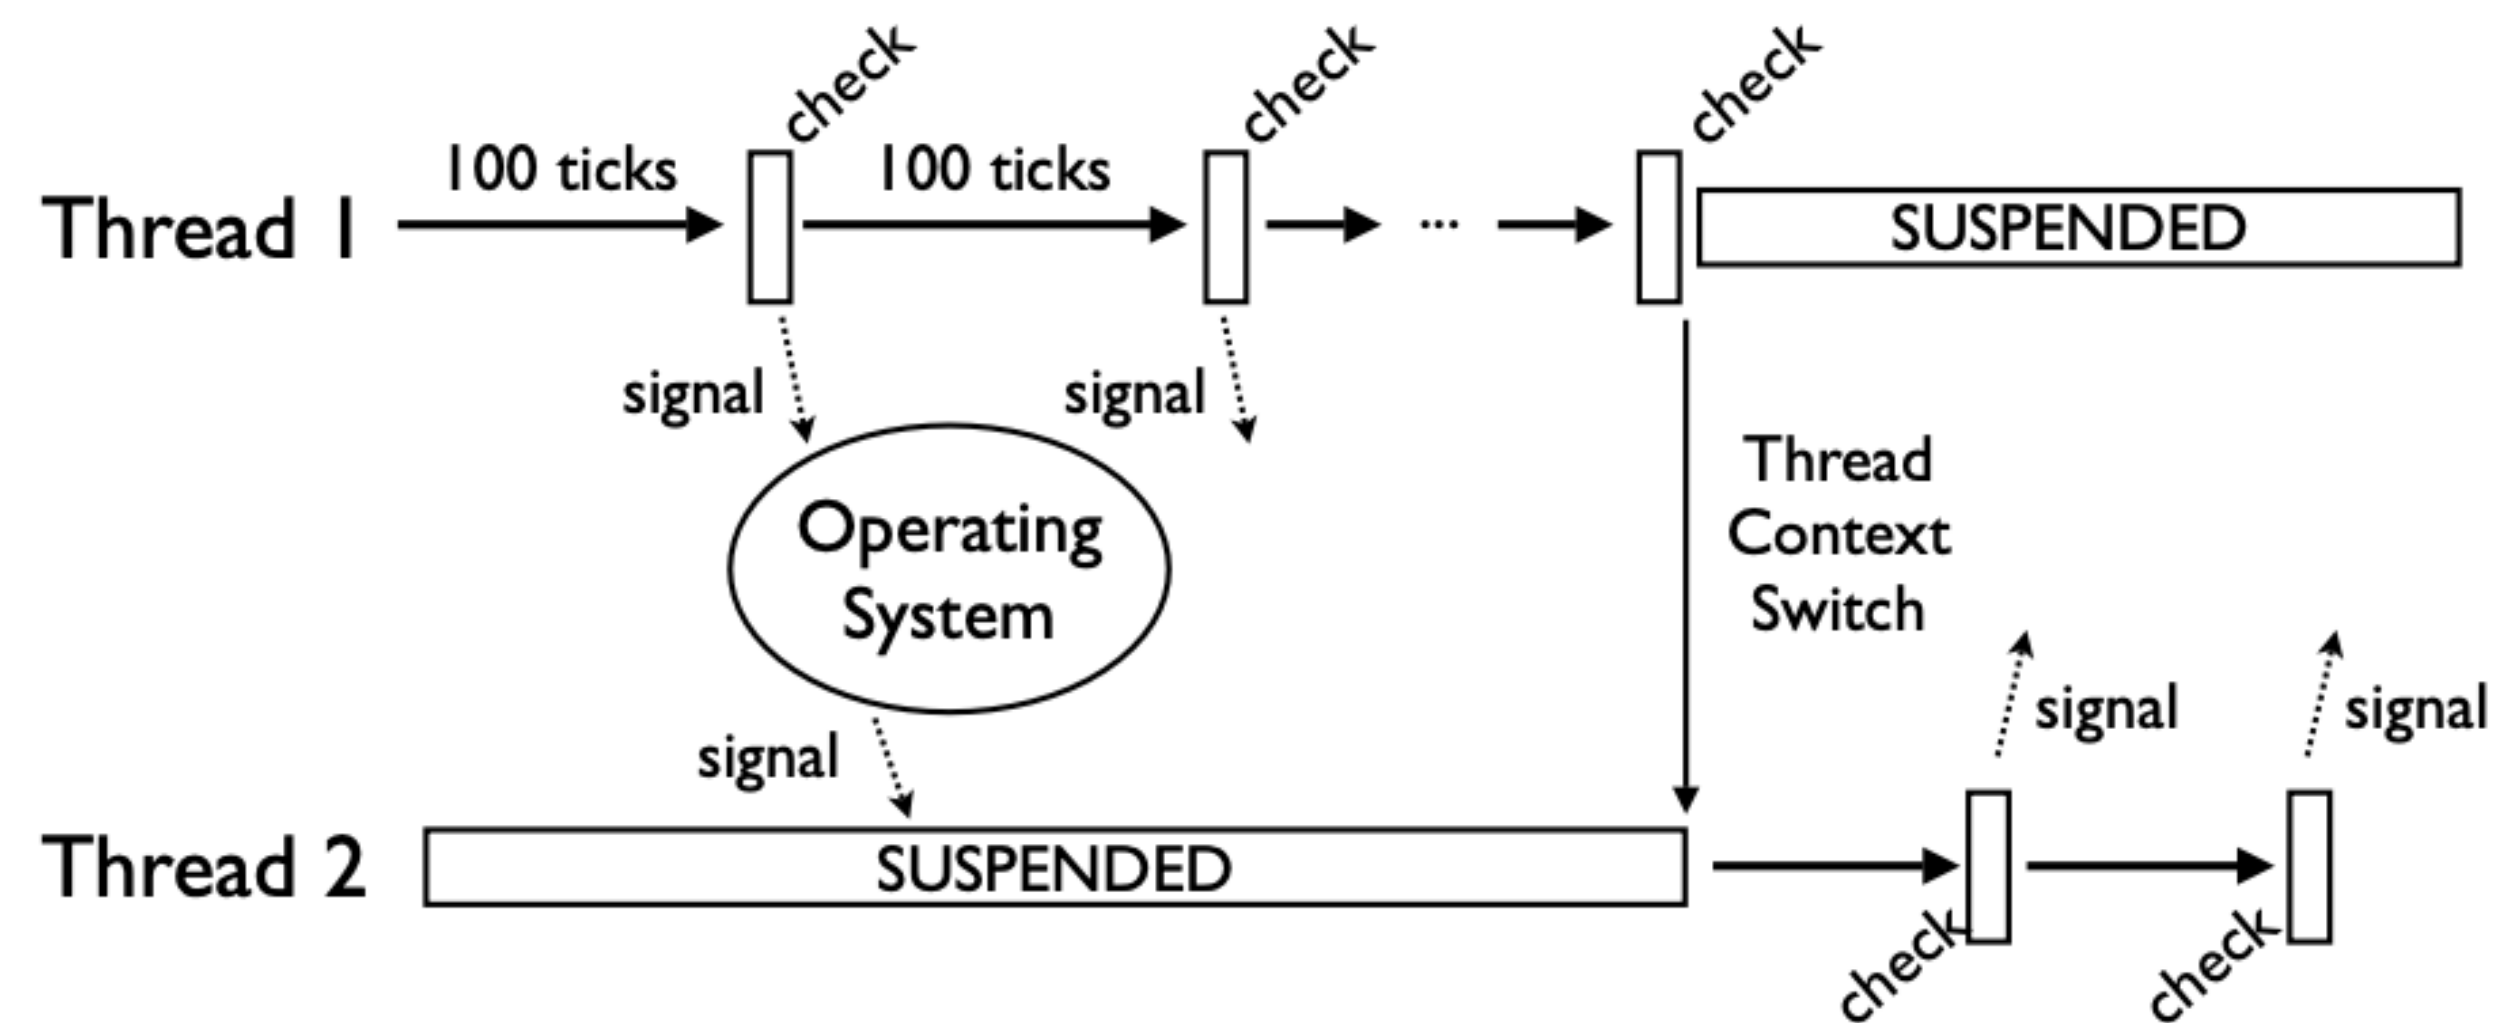
## GIL Implementation

- The GIL is not a simple mutex lock
- The implementation (Unix) is either...
  - A POSIX unnamed semaphore
  - Or a pthreads condition variable
- All interpreter locking is based on signaling
  - To acquire the GIL, check if it's free. If not, go to sleep and wait for a signal
  - To release the GIL, free it and signal

# Python and GIL Lock

## Thread Scheduling

- Thread switching is far more subtle than most programmers realize (it's tied up in the OS)



- The lag between signaling and scheduling may be significant (depends on the OS)



# Python and GIL Lock

## CPU-Bound Threads

- As we saw earlier, CPU-bound threads have horrible performance properties
- Far worse than simple sequential execution
  - 24.6 seconds (sequential)
  - 45.5 seconds (2 threads)
- A big question :Why?
  - What is the source of that overhead?

example in  
Python file  
in PyCharm

# Python and GIL Lock

## The GIL and C Code

- As mentioned, Python can talk to C/C++
- C/C++ extensions can release the interpreter lock and run independently
- Caveat : Once released, C code shouldn't do any processing related to the Python interpreter or Python objects
- The C code itself must be thread-safe

# Python and GIL Lock

## Why is the GIL there?

- Simplifies the implementation of the Python interpreter (okay, sort of a lame excuse)
- Better suited for reference counting (Python's memory management scheme)
- Simplifies the use of C/C++ extensions. Extension functions do not need to worry about thread synchronization
- And for now, it's here to stay... (although people continue to try and eliminate it)

# Processes and Messages




# Processes and Messages

## Concept: Message Passing

- An alternative to threads is to run multiple independent copies of the Python interpreter
- In separate processes
- Possibly on different machines
- Get the different interpreters to cooperate by having them send messages to each other

# Processes and Messages

## Message Passing



- On the surface, it's simple
- Each instance of Python is independent
- Programs just send and receive messages
- Two main issues
  - What is a message?
  - What is the transport mechanism?

# Processes and Messages

## Messages

- A message is just a bunch of bytes (a buffer)
- A "serialized" representation of some data
- Creating serialized data in Python is easy

# Processes and Messages

## pickle Module

- A module for serializing objects
- Serializing an object onto a "file"

```
import pickle
...
pickle.dump(someobj, f)
```

- Unserializing an object from a file

```
someobj = pickle.load(f)
```

- Here, a file might be a file, a pipe, a wrapper around a socket, etc.

# Processes and Messages

## Pickle Commentary

- Using pickle is almost too easy
- Almost any Python object works
  - Builtins (lists, dicts, tuples, etc.)
  - Instances of user-defined classes
  - Recursive data structures
- Exceptions
  - Files and network connections
  - Running generators, etc.

# Processes and Messages

## Message Transport

- Python has various low-level mechanisms
  - Pipes
  - Sockets
  - FIFOs
- Libraries provide access to other systems
  - MPI
  - XML-RPC (and many others)



# Processes and Messages


## An Example

- Launching a subprocess and hooking up the child process via a pipe
- Use the subprocess module

```
import subprocess

p = subprocess.Popen(['python', 'child.py'],
                     stdin=subprocess.PIPE,
                     stdout=subprocess.PIPE)

p.stdin.write(data)      # Send data to subprocess
p.stdout.read(size)     # Read data from subprocess
```



# Processes and Messages

## Pipes and Pickle


- Most programmers would use the subprocess module to run separate programs and collect their output (e.g., system commands)
- However, if you put a pickling layer around the files, it becomes much more interesting
- Becomes a communication channel where you can send just about any Python object



# Processes and Messages

## Big Picture

- Can easily have 10s-1000s of communicating Python interpreters



# Processes and Messages

## Interlude

- Message passing is a fairly general concept
- However, it's also kind of nebulous in Python
- No agreed upon programming interface
- Vast number of implementation options
- Intersects with distributed objects, RPC, cross-language messaging, etc.

# **Multiprocessing Module**

# Multiprocessing Module

## multiprocessing Module

- A new library module added in Python 2.6
- Originally known as pyprocessing (a third-party extension module)
- This is a module for writing concurrent Python programs based on communicating processes
- A module that is especially useful for concurrent CPU-bound processing

# Multiprocessing Module

## Using multiprocessing

- Here's the cool part...
- You already know how to use multiprocessing
- At a very high-level, it simply mirrors the thread programming interface
- Instead of "Thread" objects, you now work with "Process" objects.

# Multiprocessing Module

## multiprocessing Example

- Define tasks using a Process class

```
import time
import multiprocessing

class CountdownProcess(multiprocessing.Process):
    def __init__(self, count):
        multiprocessing.Process.__init__(self)
        self.count = count
    def run(self):
        while self.count > 0:
            print "Counting down", self.count
            self.count -= 1
            time.sleep(5)
        return
```

- You inherit from Process and redefine run()

example in  
Python file  
in PyCharm



# Multiprocessing Module

## Launching Processes

- To launch, same idea as with threads

```
if __name__ == '__main__':  
    p1 = CountdownProcess(10)    # Create the process object  
    p1.start()                  # Launch the process  
  
    p2 = CountdownProcess(20)    # Create another process  
    p2.start()                  # Launch
```

- Processes execute until run() stops
- A critical detail : Always launch in main as shown (required for Windows)

example in  
Python file  
in PyCharm

# Multiprocessing Module

## Functions as Processes

- Alternative method of launching processes

```
def countdown(count):  
    while count > 0:  
        print "Counting down", count  
        count -= 1  
        time.sleep(5)  
  
if __name__ == '__main__':  
    p1 = multiprocessing.Process(target=countdown,  
                                args=(10,))  
    p1.start()
```

- Creates a Process object, but its run() method just calls the given function

example in  
Python file  
in PyCharm



# Multiprocessing Module

## Does it Work?

- Consider this CPU-bound function

```
def count(n):  
    while n > 0:  
        n -= 1
```

- Sequential Execution:

```
count(100000000)  
count(100000000)
```

→ 24.6s

- Multiprocessing Execution

```
p1 = Process(target=count, args=(100000000,))  
p1.start()  
p2 = Process(target=count, args=(100000000,))  
p2.start()
```

→ 12.5s

- Yes, it seems to work

example in  
Python file  
in PyCharm

# Multiprocessing Module

## Other Process Features

- Joining a process (waits for termination)

```
p = Process(target=somefunc)
p.start()
...
p.join()
```

- Making a daemon process

```
p = Process(target=somefunc)
p.daemon = True
p.start()
```

- Terminating a process

```
p = Process(target=somefunc)
...
p.terminate()
```

- These mirror similar thread functions

# Multiprocessing Module

## Distributed Memory

- With multiprocessing, there are no shared data structures
- Every process is completely isolated
- Since there are no shared structures, forget about all of that locking business
- Everything is focused on messaging

**Some are implemented  
in later versions of  
Python e.g. > 3.0**

# **Multiprocessing Module - Pipes**

# Pipes in Multiprocessing Module

## Pipes

- A channel for sending/receiving objects  
`(c1, c2) = multiprocessing.Pipe()`
- Returns a pair of connection objects (one for each end-point of the pipe)
- Here are methods for communication

```
c.send(obj)           # Send an object
c.recv()              # Receive an object

c.send_bytes(buffer)  # Send a buffer of bytes
c.recv_bytes([max])   # Receive a buffer of bytes

c.poll([timeout])     # Check for data
```

# Pipes in Multiprocessing Module

## Using Pipes

- The `Pipe()` function largely mimics the behavior of Unix pipes
- However, it operates at a higher level
- It's not a low-level byte stream
- You send discrete messages which are either Python objects (pickled) or buffers



# Pipes in Multiprocessing Module

## Pipe Example

- A simple data consumer

```
def consumer(p1, p2):  
    p1.close()    # Close producer's end (not used)  
    while True:  
        try:  
            item = p2.recv()  
        except EOFError:  
            break  
        print item    # Do other useful work here
```

- A simple data producer

```
def producer(sequence, output_p):  
    for item in sequence:  
        output_p.send(item)
```

example in  
Python file  
in PyCharm

# Pipes in Multiprocessing Module

## Pipe Example

```
if __name__ == '__main__':
    p1, p2 = multiprocessing.Pipe()

    cons = multiprocessing.Process(
        target=consumer,
        args=(p1, p2))
    cons.start()

    # Close the input end in the producer
    p2.close()

    # Go produce some data
    sequence = xrange(100) # Replace with useful data
    producer(sequence, p1)

    # Close the pipe
    p1.close()
```

example in  
Python file  
in PyCharm



# **Multiprocessing Module - Queues**

# Queue in Multiprocessing Module

## Message Queues

- multiprocessing also provides a queue
- The programming interface is the same

```
from multiprocessing import Queue
```

```
q = Queue()  
q.put(item)      # Put an item on the queue  
item = q.get()   # Get an item from the queue
```

- There is also a joinable Queue

```
from multiprocessing import JoinableQueue
```

```
q = JoinableQueue()  
q.task_done()    # Signal task completion  
q.join()         # Wait for completion
```

# Queue in Multiprocessing Module

## Queue Implementation

- Queues are implemented on top of pipes
- A subtle feature of queues is that they have a "feeder thread" behind the scenes
- Putting an item on a queue returns immediately (allowing the producer to keep working)
- The feeder thread works on its own to transmit data to consumers

# Queue in Multiprocessing Module

## Queue Example

- A consumer process

```
def consumer(input_q):  
    while True:  
        # Get an item from the queue  
        item = input_q.get()  
        # Process item  
        print item  
        # Signal completion  
        input_q.task_done()
```

- A producer process

```
def producer(sequence, output_q):  
    for item in sequence:  
        # Put the item on the queue  
        output_q.put(item)
```

example in  
Python file  
in PyCharm

# Queue in Multiprocessing Module

## Queue Example

- Running the two processes

```
if __name__ == '__main__':  
    from multiprocessing import Process, JoinableQueue  
    q = JoinableQueue()  
  
    # Launch the consumer process  
    cons_p = Process(target=consumer, args=(q,))  
    cons_p.daemon = True  
    cons_p.start()  
  
    # Run the producer function on some data  
    sequence = range(100)    # Replace with useful data  
    producer(sequence, q)  
  
    # Wait for the consumer to finish  
    q.join()
```

example in  
Python file  
in PyCharm



# Queue in Multiprocessing Module

## Commentary

- If you have written threaded programs that strictly stick to the queuing model, they can probably be ported to multiprocessing
- The following restrictions apply
  - Only objects compatible with pickle can be queued
  - Tasks can not rely on any shared data other than a reference to the queue



# Queue in Multiprocessing Module

## Other Features

- multiprocessing has many other features
  - Process Pools
  - Shared objects and arrays
  - Synchronization primitives
  - Managed objects
  - Connections
- Will briefly look at one of them

# **Multiprocessing Module - Process Pools**

# Process Pools in Multiprocessing Module

## Process Pools

- Creating a process pool

```
p = multiprocessing.Pool([numprocesses])
```

- Pools provide a high-level interface for executing functions in worker processes
- Let's look at an example...

# Process Pools in Multiprocessing Module

## Pool Example

- Define a function that does some work
- Example : Compute a SHA-512 digest of a file

```
import hashlib

def compute_digest(filename):
    digest = hashlib.sha512()
    f = open(filename, 'rb')
    while True:
        chunk = f.read(8192)
        if not chunk: break
        digest.update(chunk)
    f.close()
    return digest.digest()
```

- This is just a normal function (no magic)

example in  
Python file  
in PyCharm

# Process Pools in Multiprocessing Module

## Pool Example

- Here is some code that uses our function
- Make a dict mapping filenames to digests

```
import os
TOPDIR = "/Users/beazley/Software/Python-3.0"

digest_map = {}
for path, dirs, files in os.walk(TOPDIR):
    for name in files:
        fullname = os.path.join(path, name)
        digest_map[fullname] = compute_digest(fullname)
```

- Running this takes about 10s on my machine

example in  
Python file  
in PyCharm



# Process Pools in Multiprocessing Module

## Pool Example

- With a pool, you can farm out work
- Here's a small sample

```
p = multiprocessing.Pool(2)    # 2 processes

result = p.apply_async(compute_digest, ('README.txt',))
...
... various other processing
...
digest = result.get()          # Get the result
```

- This executes a function in a worker process and retrieves the result at a later time
- The worker churns in the background allowing the main program to do other things

example in  
Python file  
in PyCharm



# Process Pools in Multiprocessing Module

## Pool Example

- Make a dictionary mapping names to digests

```
import multiprocessing
import os
TOPDIR = "/Users/beazley/Software/Python-3.0"

p = multiprocessing.Pool(2)      # Make a process pool
digest_map = {}
for path, dirs, files in os.walk(TOPDIR):
    for name in files:
        fullname = os.path.join(path, name)
        digest_map[fullname] = p.apply_async(
            compute_digest, (fullname,)
        )

# Go through the final dictionary and collect results
for filename, result in digest_map.items():
    digest_map[filename] = result.get()
```

- This runs in about 5.6 seconds

example in  
Python file  
in PyCharm

# Resources

- **An Introduction to Python Concurrency.** <http://www.dabeaz.com/userix2009/concurrent/index.html>
- **Slides on Concurrency** <http://www.dabeaz.com/userix2009/concurrent/Concurrent.pdf>
- **Python Multiprocessing: The Complete Guide** <https://superfastpython.com/multiprocessing-in-python/>

**The End**